

# DEEP LEARNING ON GRAPHS

## LEARNING BEYOND EUCLIDEAN DATA

---

Michaël DEFFERRARD

École Polytechnique Fédérale de Lausanne (EPFL)

Joint work with    Youngjoo SEO (EPFL)  
                      Xavier BRESSON (NTU)  
                      Pierre VANDERGHEYNST (EPFL)

# Structured data

## Why structure data ?

- ▶ To incorporate additional information.
- ▶ To exploit spatial correlations.
- ▶ To decrease learning complexity by making geometric assumptions.

## Data structured by Euclidean grids.

- ▶ 1D: sound, time-series.
- ▶ 2D: images.
- ▶ 3D: video, hyper-spectral images.

# Naturally graph-structured data


Modeling versatility: graphs model heterogeneous pairwise relationships.

Examples of irregular / graph-structured data:

- ▶ Social networks: Facebook, Twitter.
- ▶ Biological networks: genes, molecules, brain connectivity.
- ▶ Infrastructure networks: energy, transportation, Internet, telephony.



Social network



Brain structure



Telecommunication

# Types of graphs


## Sample graph

- ▶ Semi-supervised learning.
- ▶ Incorporate external information.

## Feature graph

- ▶ Reduce computations.
- ▶ Incorporate external information.

 document  
 citation  
 hyper-link  
 words



 word  
 similarity  
 documents

Problems: signals, nodes or graphs classification (regression).

## Using the structure

Extrinsic: embed the graph in an Euclidean space.

- ▶ Each node is represented by a vector.
- ▶ Use that embedding as additional features for a fully connected NN.
- ▶ Use a convolutional NN in the embedding space.  
**Possibly very high-dimensional!**

Intrinsic: a Neural Net defined on graphically structured data.

- ▶ Exploit geometric structure for computational efficiency.
- ▶ Starting point: ConvNet, an intrinsic formulation for Euclidean grids.

# Why ConvNets?

ConvNets are extremely efficient at extracting meaningful statistical patterns in large-scale and high-dimensional datasets.

They exploit the geometry.

## Statistical assumptions

- ▶ Localization: compact filters for low complexity.
- ▶ Stationarity: translation invariance.
- ▶ Compositionality: analysis with a filterbank.
- ▶ Multi-scale: hierarchical features extracted by multiple layers.


# Developed for data lying on Euclidean grids

All operations are well defined and computationally efficient:

1. Convolution → filter translation or fast Fourier transform (FFT).
2. Down-sampling → pick one pixel out of  $n$ .
3. Non-linearity → point-wise operation.
4. Pooling → summarize the receptive field.



Image (2D)   Video (3D)





Sound (1D)

# ConvNets on graphs

## Graphs vs Euclidean grids

- ▶ Irregular sampling.
- ▶ Weighted edges.
- ▶ No orientation (in general).




## Challenges

1. Formulate convolution and down-sampling on graphs.
2. Make them efficient!

# ConvNets on graphs: spatial approach


Niepert, Ahmed, and Kutzkov 2016

1. Define receptive field / neighborhood.
2. Order nodes, i.e. give orientations to a node's edges.




# ConvNets on graphs: spectral approach

Bruna, Zaremba, Szlam, and LeCun 2014; Henaff, Bruna, and LeCun 2015



# Graphs

$\mathcal{G} = (\mathcal{V}, \mathcal{E}, W)$ : undirected and connected graph



- ▶  $\mathcal{V}$ : set of  $|\mathcal{V}| = n$  vertices
- ▶  $\mathcal{E}$ : set of edges
- ▶  $W \in \mathbb{R}^{n \times n}$ : weighted adjacency matrix
- ▶  $D_{ii} = \sum_j W_{ij}$ : diagonal degree matrix

Graph Laplacians (core operator to spectral graph theory):

- ▶  $L = D - W \in \mathbb{R}^{n \times n}$ : combinatorial
- ▶  $L = I_n - D^{-1/2} W D^{-1/2}$ : normalized


# Graph Fourier basis

Shuman, Narang, Frossard, Ortega, and Vandergheynst 2013

$L$  is symmetric and positive semidefinite  $\rightarrow L = U \Lambda U^T$  (EVD)

- ▶ Graph Fourier basis  $U = [u_1, \dots, u_n] \in \mathbb{R}^{n \times n}$


- ▶ Graph “frequencies”  $\Lambda = \begin{bmatrix} \lambda_1 & & 0 \\ & \ddots & \\ 0 & & \lambda_n \end{bmatrix} \in \mathbb{R}^{n \times n}$



# Graph Fourier Transform

Shuman, Narang, Frossard, Ortega, and Vandergheynst 2013

- ▶ Graph signal  $x : \mathcal{V} \rightarrow \mathbb{R}$  seen as  $x \in \mathbb{R}^n$
- ▶ Transform:  $\hat{x} = \mathcal{F}_{\mathcal{G}}\{x\} = U^T x \in \mathbb{R}^n$
- ▶ Inverse:  $x = \mathcal{F}_{\mathcal{G}}^{-1}\{x\} = U\hat{x} = UU^T x = x$




# Filtering with convolution on graphs

Shuman, Narang, Frossard, Ortega, and Vandergheynst 2013


$$\text{Convolution theorem: } y = x *_{\mathcal{G}} g = U (U^T g \odot U^T x) = U (\hat{g} \odot U^T x)$$

$$y = x *_{\mathcal{G}} g = U \begin{bmatrix} \hat{g}(\lambda_0) & & 0 \\ & \ddots & \\ 0 & & \hat{g}(\lambda_{n-1}) \end{bmatrix} U^T x = U \hat{g}(\Lambda) U^T x = \hat{g}(L)x$$


$y = \hat{g}(L)\delta_{10}$ : localized on sensor



$$y^T y = 367.18$$




$$y^T y = 6.83$$




$$y^T y = 0.00$$


$\hat{g}(\lambda)$ : filter defined in the spectral domain



$$\hat{g}(\lambda) = \exp\left(-\frac{0\lambda}{\lambda_{\max}}\right)$$






$$\hat{g}(\lambda) = \exp\left(-\frac{5\lambda}{\lambda_{\max}}\right)$$



$$\hat{g}(\lambda) = \exp\left(-\frac{100\lambda}{\lambda_{\max}}\right)$$


$y = \hat{g}(L)\delta_{50}$ : localized on ring graph



# Spectral filtering of graph signals

Non-parametric filter, can learn all possible filters:

$$\hat{g}_\theta(\Lambda) = \text{diag}(\theta), \quad \theta \in \mathbb{R}^n$$



- ▶ Non-localized in vertex domain
- ▶ Learning complexity in  $\mathcal{O}(n)$
- ▶ Computational complexity in  $\mathcal{O}(n^2)$  (& memory)

Variation: a smooth function such as  $\hat{g}_\theta(\Lambda) = B\theta$  where  $B$  is the cubic spline basis (Bruna, Zaremba, Szlam, and LeCun 2014).

# Polynomial parametrization

Shuman, Ricaud, and Vandergheynst 2016


$$\hat{g}_\theta(\Lambda) = \sum_{k=0}^{K-1} \theta_k \Lambda^k, \quad \theta \in \mathbb{R}^K$$

- ▶ Can learn all  $K$ -localized filters.
- ▶ Distributed computing: only need access to the  $K$ -neighborhood.
  
- ▶  $K$ -localized
- ▶ Learning complexity in  $\mathcal{O}(K)$
- ▶ Computational complexity in  $\mathcal{O}(n^2)$

# Filter localization

Hammond, Vandergheynst, and Gribonval 2011, Lemma 5.2

- ▶ Value at  $j$  of  $g_\theta$  centered at  $i$ :  $(\hat{g}_\theta(L)\delta_i)_j = (\hat{g}_\theta(L))_{i,j} = \sum_k \theta_k(L^k)_{i,j}$
- ▶  $d_G(i,j) > K$  implies  $(L^K)_{i,j} = 0$



# Filter localization

Shuman, Ricaud, and Vandergheynst 2016




Figure: Localization on regular Euclidean grid.





Figure: Localization on graph with  $(\hat{g}_\theta(L)\delta_i)_j = (\hat{g}_\theta(L))_{i,j}$ .

# Recursive formulation with Chebyshev polynomials

Hammond, Vandergheynst, and Gribonval 2011

$$\hat{g}_\theta(\Lambda) = \sum_{k=0}^{K-1} \theta_k T_k(\tilde{\Lambda}), \quad \tilde{\Lambda} = 2\lambda_n^{-1}\Lambda - I_n$$

Chebyshev polynomials:  $T_k(x) = 2xT_{k-1}(x) - T_{k-2}(x)$   
with  $T_0 = 1$  and  $T_1 = x$



## Recursive formulation with Chebyshev polynomials

$$y = \hat{g}_\theta(L)x = \sum_{k=0}^{K-1} \theta_k T_k(\tilde{L})x, \quad \tilde{L} = 2\lambda_n^{-1}L - I_n$$

Recurrence:

$$\begin{aligned} y &= \hat{g}_\theta(L)x = [\bar{x}_0, \dots, \bar{x}_{K-1}] \theta \\ \bar{x}_k &= T_k(\tilde{L})x = 2\tilde{L}\bar{x}_{k-1} - \bar{x}_{k-2} \\ \bar{x}_0 &= x \\ \bar{x}_1 &= \tilde{L}x \end{aligned}$$

- ▶  $K$ -localized
- ▶ Learning complexity in  $\mathcal{O}(K)$
- ▶ Computational complexity in  $\mathcal{O}(K|\mathcal{E}|)$  (same as classical ConvNets!)

# Learning filters

Defferrard, Bresson, and Vandergheynst 2016

$$y_{s,j} = \sum_{i=1}^{F_{in}} \hat{g}_{\theta_{i,j}}(L) x_{s,i} \in \mathbb{R}^n$$

- ▶  $x_{s,i}$ : feature map  $i$  of sample  $s$
- ▶  $\theta_{i,j}$ : trainable parameters  
( $F_{in} \cdot F_{out}$  vectors of  $K$  Chebyshev coefficients)




Gradients for backpropagation:

- ▶  $\frac{\partial E}{\partial \theta_{i,j}} = \sum_{s=1}^S [\bar{x}_{s,i,0}, \dots, \bar{x}_{s,i,K-1}]^T \frac{\partial E}{\partial y_{s,j}}$
- ▶  $\frac{\partial E}{\partial x_{s,i}} = \sum_{j=1}^{F_{out}} g_{\theta_{i,j}}(L) \frac{\partial E}{\partial y_{s,j}}$

Overall cost of  $\mathcal{O}(K|\mathcal{E}|F_{in}F_{out}S)$  operations,  $|\mathcal{E}| \propto n$  for sparse graphs

# Coarsening



Defferrard, Bresson, and Vandergheynst 2016



- ▶ Inherently combinatorial problem.
- ▶ Can be done as pre-processing.
- ▶ Greedy node merging with Graclus / Metis (very fast).

# Pooling

Defferrard, Bresson, and Vandergheynst 2016




Pooling as any regular 1D signal

- ▶ Node order does not matter → arrange them for local access.
- ▶ Nodes at multiple levels are ordered as a tree.
- ▶ Satisfies parallel architectures like GPUs.

# Graph ConvNet architecture

Defferrard, Bresson, and Vandergheynst 2016



# Revisiting MNIST

Model	Architecture	Accuracy
Classical CNN	C32-P4-C64-P4-FC512	99.33
Proposed graph CNN	GC32-P4-GC64-P4-FC512	99.14

Table: Comparison to classical ConvNets.


Comparable to classical ConvNets,  
and better than other parametrizations!

Architecture	Accuracy		
	Non-Param	Spline	Chebyshev
GC10	95.75	97.26	97.48
GC32-P4-GC64-P4-FC512	96.28	97.15	99.14

Table: Comparison between spectral filters,  $K = 25$ .

# 20NEWS: structuring documents with a feature graph

Defferrard, Bresson, and Vandergheynst 2016



# Classification accuracy & Graph quality

Model	Accuracy
Linear SVM	65.90
Multinomial Naive Bayes	68.51
Softmax	66.28
FC2500	64.64
FC2500-FC500	65.76
GC32	68.26

Proposed graph CNN vs. other methods on 20NEWS.

bag-of-words	word2vec			
	pre-learned	learned	approximate	random
67.50	66.98	68.26	67.86	67.75


Different graphs on 20NEWS.

Architecture	8-NN on 2D Euclidean grid	random
GC32	97.40	96.88
GC32-P4-GC64-P4-FC512	99.14	95.39

Different graphs on MNIST.

# 20NEWS: training time

Defferrard, Bresson, and Vandergheynst 2016



Make CNNs practical for graph signals!

Spline:  $\hat{g}_\theta(\Lambda) = B\theta$  where  $B$  is the cubic spline basis  
(Bruna, Zaremba, Szlam, and LeCun 2014)

# Recurrent Graph Convolutional Network

Seo, Defferrard, Bresson, and Vandergheynst 2016

## 1D signals

- ▶  $h_t = \tanh(W_x x_t + W_h h_{t-1})$
- ▶  $y_t = Wh_t$
- ▶ State stored in hidden units

## Graph signals

- ▶  $h_t = \tanh(W_x *_{\mathcal{G}} x_t + W_h *_{\mathcal{G}} h_{t-1})$
- ▶  $y_t = W *_{\mathcal{G}} h_t$
- ▶ State stored locally on the nodes

- ▶ Graph filtering  $x$  as  $y = [\bar{x}_0, \dots, \bar{x}_{K-1}]^T$  is a weighted sum of diffused versions  $\bar{x}$  of  $x$ .
- ▶ Data exchanged locally around the  $K$ -neighborhood.
- ▶ Reduces to independent signals if  $K = 1$  or graph has no edge.

# Various applications

- ▶ Semi-supervised learning  
[Kipf and Welling 2016; Manessi, Rozza, and Manzo 2017]
- ▶ Quantum Chemistry  
[Duvenaud et al. 2015; Gilmer, Schoenholz, Riley, Vinyals, and Dahl 2017]
- ▶ High Energy Physics
- ▶ Computer Graphics [Monti, Boscaini, et al. 2016; Yi, Su, Guo, and Guibas 2016; Wang, Gan, Zhang, and Shui 2017; Simonovsky and Komodakis 2017]
- ▶ Community detection [Bruna and Li 2017]
- ▶ Brain analysis  
[Ktena et al. 2017; Parisot et al. 2017; Anirudh and Thiagarajan 2017]
- ▶ Matrix completion for recommendation  
[Monti, Bronstein, and Bresson 2017]
- ▶ Neural machine translation  
[Bastings, Titov, Aziz, Marcheggiani, and Sima'an 2017]
- ▶ Link prediction and entity classification in knowledge bases  
[Schlichtkrull et al. 2017]

# Conclusion

- ▶ Many problems can be defined as data structured by graphs.
- ▶ ConvNets are efficient because they exploit the geometry.
- ▶ Spatial approaches require nodes to be ordered.
- ▶ Convolution is well-defined in the Fourier domain, though the graph Fourier transform (GFT) is expensive.
- ▶ Polynomial formulations of order  $K$  are  $K$ -localized and do not require the GFT.
- ▶ Graph coarsening is a research problem in itself.
- ▶ Pooling made efficient by ordering nodes and inserting fake nodes.
- ▶ Loss of edge orientation is not killing performance on images.
- ▶ Having a good (relative to the task) graph is important.
- ▶ RNN extension can accommodate structured time series.
- ▶ Many proposed applications and many more to come.

# Thanks Questions?

- ▶ **Paper:** Defferrard, Bresson and Vandergheynst, Convolutional Neural Networks on Graphs with Fast Localized Spectral Filtering, NIPS, 2016.
- ▶ **Code:** [https://github.com/mdeff/cnn\\_graph](https://github.com/mdeff/cnn_graph)
- ▶ **Paper:** Seo, Defferrard, Bresson and Vandergheynst, Structured Sequence Modeling with Graph Convolutional Recurrent Networks, arXiv, 2017.
- ▶ **Code:** <https://github.com/youngjoo-epfl/gconvRNN>